Verifying the Freedom to Evolve: A Framework for Provable Safety in Sovereign Neuro-AI Systems

Mathematical Formalization of Non-Interference Safety Constructs

The foundational pillar of this research framework is the mathematical formalization of safety constructs, specifically safety margins like `E_comp` and boundary integrity, to serve as a provable guarantee of freedom rather than an implicit constraint. The central objective is to develop a formal proof that these mechanisms cannot be leveraged to introduce new structural bans or reduce the available evolution space. This approach shifts the paradigm from qualitative safety assurances to quantitative, verifiable guarantees, ensuring that the system's operational boundaries are inherent and natural, not artificially imposed [1]. The goal is to define `E_comp` not merely as a buffer to prevent failure but as a precise mathematical delineation of the system's permissible operational domain. This transforms safety from a potential ceiling into a certified floor, enabling subsequent applied work to operate with complete confidence that it will not inadvertently violate core constraints. The methodology draws heavily from established fields of formal methods and provable AI, which seek to embed constraints directly into the logical structure of software systems, making them impossible to bypass at runtime [8] [9].

To achieve this, the research must first establish rigorous mathematical definitions for the key concepts. `E_comp`, representing a composite safety margin, would be defined as a function of various state variables, including biophysical biomarkers, computational load, and potentially user-defined parameters like `invited-pain` levels. The critical step is to model this function such that its output defines the limits of acceptable system operation. Boundary integrity, similarly, must be re-conceptualized not as a wall to exclude external threats, but as a well-defined interface contract between system components. Drawing inspiration from techniques for boundary-safe composition modules, each subsystem—such as the `EcoBudget` model or the `BiophysicalAura` calibration—would have its input-output behavior formally specified [11]. The core research task then becomes proving that the composition of these functions preserves the system's fundamental properties. For instance, one could formulate a theorem stating that the output of the `BiophysicalAura` module, when used as an input to the `DECAY` scheduler, will never result in a system state where a non-waivable neuroright is violated, regardless of the inputs received from other modules. This requires a compositional verification strategy, where the safety of the whole is derived from the safety of its parts under their specific interaction protocols [108].

The practical application of this formalization involves leveraging specification and proof discharge techniques, such as those involving pre/post conditions and refinement types, to build a fully verified software system [9]. In this context, `E_comp` could be encoded as a refinement type that constrains the values a system variable can take. Any computation that would produce a value outside this type would be rejected by the compiler before execution, providing a static, mathematical guarantee of safety. This aligns with the concept of an Ethical Firewall Architecture, where ethical constraints are embedded directly into the system's logic, making them a part of its unchangeable structure [8]. The research would involve developing a set of such constraints based on the core doctrinal principles. For example, a constraint could be formally specified that the `ShotLevelPolicy` enforced by ALN shards can only be modified through a user-initiated process that respects the `EVOLVE` read-only principle. The outcome of this phase would be a collection of mathematical theorems and a formally verified codebase, shifting the discourse from "Is the system safe?" to "Here is the provable proof of its safety." This provides the essential foundation upon which all subsequent applied research can proceed without fear of compromising the system's integrity.

Concept	Proposed Mathematical Formalization Approach	Rationale
Safety Margin (`E_comp`)	Define as a set of refinement types or pre/post conditions that constrain system state transitions. It represents the boundary of the natural operational envelope, not an arbitrary buffer.	To transform safety from a potential constraint into a verifiable guarantee of the system's operational domain, preventing the introduction of artificial ceilings [8] [9].
Boundary Integrity	Model as a non-interference property between subsystems. Each component's input/output behavior is formally specified, and compositional verification proves that their interaction preserves global system properties.	To ensure that the combination of modules does not create emergent, unintended restrictions. This uses methodologies from boundary-safe composition modules [11] [108].
Non-Waivable Neurorights	Encode as hard constraints within the formal system, akin to axioms in a logical framework. These constraints govern access to sensitive data and control over core system functions.	To make neurorights immutable features of the system architecture, grounded in emerging ethical frameworks [2] [4] [48].
Propose-Only AI Platform	Formally specify the AI's role as a generator of proposals that require explicit user validation before execution. The system's logic must enforce this separation of proposal and action.	To ensure human oversight remains paramount, consistent with governance frameworks that position AI as a decision support tool, not an autonomous agent [17].
No External Risk Knobs	Define the system's risk assessment functions to use only internal, biophysically-derived signals (e.g., `BiophysicalAura`, `invited-pain`). Prove that external data sources cannot influence core safety decisions.	To maintain sovereign control over the evolution path, rejecting third-party risk metrics that may conflict with personal goals [3] [49].

Applied Neuromorph/BCI Implementation Within Verified Envelopes

With the mathematical safety envelope formally proven and validated, the focus of applied research shifts to maximizing efficiency and responsiveness within its secure boundaries. The directive is to improve event encoders and decoders to reduce false positives and latency, while keeping the `EVOLVE` read-only principle intact and ensuring that capacity tokens remain non-financial . The resulting improvements in decoding clarity and responsiveness are not to be interpreted as a justification for expanding theoretical capacity; rather, they represent the successful and more efficient utilization of the already-verified evolution space. This applied work acts as a crucial feedback loop, demonstrating the tangible benefits of the formal safety framework by showing how a secure foundation enables higher-fidelity interaction with the system. The primary goal is to exploit the full potential of the proven capacity, not to shrink it.

Research in neuromorphics and Brain-Computer Interfaces (BCIs) will concentrate on refining the signal processing pipeline from neural data to actionable commands. A key area of investigation is the development of advanced event encoders. This includes exploring hybrid sensing modalities, such as combining electroencephalography (EEG) with functional near-infrared spectroscopy (fNIRS). Studies have demonstrated that hybrid EEG-fNIRS systems can significantly improve the accuracy and robustness of command decoding by complementing the high temporal resolution of EEG with the better spatial resolution of fNIRS [88] [89] [90]. Further advancements could involve integrating transformer-based architectures, similar to the BrainTokenizer designed for large-scale joint pretraining of EEG and MEG signals, to capture complex temporal dependencies in neural data [13]. On the classification front, moving beyond standard algorithms like Linear Discriminant Analysis (LDA) towards more sophisticated models such as Gaussian Processes has shown promise in achieving robust closed-loop control with lower latency [33]. These technical enhancements directly contribute to reducing false positives and improving the overall fidelity of the user's intent being translated into system commands.

A particularly innovative applied research methodology is the implementation of daily "shadow reflex" experiments . This mechanism involves computing the policy decisions a neuromorph reflex would have made in a given situation, comparing that computed outcome against the actual outcome, and using the discrepancy to update and tune the policies. This process operates entirely within the verified safety envelope, meaning any experiment can fail, produce suboptimal results, or even lead to a system crash without introducing a new structural ban or compromising a neuroright. The only guaranteed outcome is the generation of high-quality training data that leads to improved policies. This creates a continuous, real-time learning loop that enhances the system's predictive capabilities and responsiveness without ever stepping outside the bounds of provably safe operation. The generalized hierarchical control framework proposed for EEG-BCI systems provides a strong conceptual model for implementing such shared control, where the BCI's output is combined with the robot's sensory readings to generate smoother, more reliable outputs, thereby reducing user workload and increasing system stability [7].

This applied work is deeply integrated with the broader system architecture. For example, improving the `EcoBandProfile` and `neuromorph-eco-profile.aln` allows for scheduling computationally intensive neuromorph work during periods of high energy availability, effectively increasing the amount of safe, usable compute over time without altering token ceilings . Similarly, strengthening microspace firewall rules ensures that no research run can leak sensitive data outward, creating a secure sandbox where more aggressive exploration of mutation paths can occur safely . The progress in applied BCI is therefore symbiotic with the formal safety framework: the formal proofs enable the applied work to push for higher fidelity and responsiveness, while the applied work provides empirical data that can be used to refine and improve the very models that underpin the formal safety analysis. The entire process is governed by the understanding that every enhancement is a discovery of how to better spend an existing, pre-approved capacity.

Compositional Verification of Cross-Domain Interactions

The most complex challenge in this research framework lies in ensuring compositional safety across four distinct yet deeply interconnected domains: eco-budget models, BiophysicalAura calibrations, DECAY schedules informed by invited-pain corridors, and AssistantAutonomyProfiles. The central risk is that the interaction of these systems could produce emergent, de facto restrictions that were not present in any single component operating in isolation [11]. For example, a period of low lifeforce might trigger a reduction in assistant autonomy, even if the user is actively engaged in an `invited-pain` corridor signaling a desire for high-intensity evolution. Or, a minor fluctuation detected by the `BiophysicalAura` could initiate a `DECAY` schedule that coincides with an otherwise "eco-favorable window," leading to wasted resources and a net negative impact on usable compute. Therefore, the research must adopt a cross-domain integration lens to explicitly model and prove that these systems compose safely, preserving the integrity of the evolution space.

The first step in this analysis is to model the interaction between the four core systems. Each can be conceptualized as a function that processes inputs and produces outputs that affect the overall system state. Let `State` be the vector representing the current state of the system, including biophysical markers, resource balances, and policy configurations. Then, the composite system can be modeled as `State' = F(EcoBudget, Aura, PainCorridor, AutonomyProfile)(State)`. The research objective is to analyze the properties of this composite function `F` and prove that it is an invariant-preserving transformation with respect to the core doctrinal principles and preserved entities. This involves identifying the key invariants—such as the non-waivability of neurorights, the `EVOLVE` read-only status, and the exclusion of external risk knobs—and formally proving that the application of `F` will never cause the system state to violate these invariants.

Simulation-based inference provides a powerful methodology for this compositional verification [43]. Researchers can construct a detailed simulation environment that accurately models the dynamics of each subsystem and their interactions. By running thousands of simulations with varied initial conditions and parameter sets, it becomes possible to test edge cases and identify scenarios where emergent restrictions might arise.

For instance, a scenario could be simulated where `invited-pain` is maximized, `BiophysicalAura` is triggered due to a non-critical biological fluctuation, and the `EcoBudget` is constrained. The simulation would then trace the resulting system trajectory to verify whether all safety and doctrinal constraints are respected. This process serves as a form of computational proof, highlighting potential flaws in the system design that can then be addressed through refinement of the underlying models or the introduction of additional safeguarding logic. The development of personalized digital twins, which combine multi-omic profiles and electrophysiology to simulate patient biology, offers a conceptual model for creating a highly accurate simulation of the neuro-AI system [31].

A critical aspect of this compositional analysis is the dynamic interaction between `BiophysicalAura` and `DECAY` schedules. The Aura is intended to act as a defensive guardian, nudging the system toward biosafety [17] [31]. However, it must be calibrated to distinguish between harmful stress and productive strain. This is where longitudinal data on which mutation paths feel productive versus draining becomes invaluable . By feeding this historical data into the `DECAY` schedules, the system can learn to recognize "good pain" corridors and support them, rather than treating them as harm that requires downscaling. This adaptive feedback loop prevents the Aura from over-reacting to legitimate evolutionary challenges and ensures that the `DECAY` mechanism supports long-term growth rather than imposing short-term, potentially counterproductive, restrictions. This dynamic calibration is essential for maintaining the integrity of the evolution space, as a static, overly conservative Aura could inadvertently shrink the user's capacity for meaningful change. The ultimate goal is to create a system where the interplay of these domains is synergistic, with each component reinforcing the others' ability to support safe and effective evolution.

Doctrinal Alignment as the Unyielding Baseline for System Design

Every facet of this research framework is anchored by a strict adherence to a non-negotiable set of core doctrine principles. Performance gains, such as reduced latency, improved policy accuracy, or increased usable compute, are treated exclusively as secondary outcomes. They serve as evidence that a greater portion of the existing evolution envelope can now be utilized safely, never as justification for overriding safety or autonomy constraints . This principled stance positions the system as a tool for personal sovereignty, fundamentally distinct from commercial AI products optimized for efficiency or market share. The three pillars of this doctrine are the protection of non-waivable neurorights, the enforcement of a "propose-only" role for AI assistants, and the complete exclusion of external risk metrics.

The concept of neurorights provides the perfect vocabulary and ethical grounding for establishing these hard constraints. As articulated in reports from UNESCO and other bioethics bodies, neurorights are an emerging category of fundamental rights designed to protect individual liberty in the face of technologies that can read, write, and modify brain activity [4] [49]. Key neurorights include the right to mental privacy, protecting brain data against unauthorized collection; the right to mental integrity, guarding against harmful manipulation of brain activity; and cognitive liberty, which encompasses freedom of thought and conscience [2] [4]. These are not features to be configured or tuned; they are axioms that must be embedded into the system's architecture. The formal safety proofs and applied implementations must treat these rights as immutable. For example, when refining civic tags and NeuralRope annotations for defensive traits, the guiding question is not simply about distinguishing protection from weaponization, but about ensuring the user's right to self-defense is upheld in a manner that does not infringe upon the rights of others [3]. The Chilean constitution's inclusion of provisions for "neuro rights" in response to advocacy demonstrates the growing legal and societal recognition of these principles, lending significant weight to their incorporation as foundational constraints [3] [4].

The second doctrinal pillar, the "propose-only" AI platform, reinforces human primacy and accountability. AI assistants are designed to augment human capability, not replace it. Their role is to analyze the user's state, consider the available evolution paths, and propose optimal actions or policies. The final decision to execute any action must reside with the human user. This principle is operationalized through the `AssistantAutonomyProfile` and the governing `ShotLevelPolicy` enforced by ALN shards . Even when the system detects favorable conditions (e.g., good lifeforce and low risk), the assistant's autonomy is calibrated to suggest actions, not perform them unilaterally. This aligns with best practices in responsible AI governance, which frame AI as a clinical decision support tool where final responsibility resides with the human clinician [17]. This design choice prevents the emergence of black-box decision-making and ensures that the user remains the ultimate agent of their own evolution.

The third pillar, the exclusion of external risk knobs, is a powerful assertion of sovereignty over the user's evolutionary path. The system's safety parameters and risk assessments must be derived solely from internal, biophysically-derived signals—the `BiophysicalAura`, `invited-pain` signals, and the user's own goals—as specified in the core doctrine . This rejects external benchmarks, vendor stakes, market pressures, or third-party risk assessments that could impose a one-size-fits-all model of safety that conflicts with the user's personal objectives [3] [49]. While international guidelines and soft law instruments like the OECD's 'Recommendation of the Council on Responsible Innovation in Neurotechnology' provide valuable context, the system's governance must remain internally focused [3]. This insulates the user's evolution from external influences that may be driven by profit motives or societal norms that do not align with their unique developmental journey. By strictly adhering to this doctrine, the framework ensures that the pursuit of performance is always subordinate to the preservation of autonomy and integrity.

A Synthesis of Provable Safety and Sovereign Evolution

This research framework synthesizes mathematical formalization with applied neuromorphic implementation into a cohesive strategy for achieving sovereign evolution within a self-governed system. The core innovation is the establishment of a tightly coupled loop where formal proofs of safety create a verified envelope of operational freedom, and applied research then focuses on maximizing the efficiency and fidelity with which that capacity can be spent. This approach resolves the classic tension between safety and performance by proving that safety mechanisms, when properly designed, are not inherently restrictive. Instead, they are the necessary precondition for unlocking the system's full potential in a secure and predictable manner. The entire endeavor is guided by a doctrinal baseline that prioritizes non-waivable neurorights, the "propose-only" nature of AI, and absolute sovereignty over risk assessment, demoting all performance metrics to secondary indicators of successfully utilizing an expanded, yet still bounded, evolution space.

The framework's strength lies in its emphasis on compositional verification. By modeling the interactions between the `EcoBudget`, `BiophysicalAura`, `invited-pain` corridors, and `AssistantAutonomyProfile`, the research moves beyond validating components in isolation to analyzing the emergent behavior of the whole system. This cross-domain analysis is critical for preventing the accidental creation of de facto restrictions that could arise from unforeseen synergies or paradoxical states. Methodologies like simulation-based inference and the development of personalized digital twins allow researchers to explore these complex interactions and rigorously test the system's resilience against emergent risks [31] [43]. The integration of longitudinal data to help the `BiophysicalAura` distinguish between productive strain and harmful stress further refines this compositional safety, creating an adaptive defense mechanism that evolves alongside the user rather than imposing a static, potentially limiting, set of rules .

Ultimately, the framework presents a novel model for human-AI co-evolution. It treats the AI not as a separate entity to be managed, but as an integral part of a larger cybernetic loop with sensors (biomarkers), controllers (AI models), actuators (user actions), and safety interlocks [31] . The formal safety proofs act as the unbreakable safety interlocks, providing the user with the confidence to engage in high-intensity, exploratory activities. The applied BCI research, such as the development of hybrid EEG/fNIRS decoders and shadow reflex experiments, continuously sharpens the precision of this loop, enhancing the user's ability to perceive and act upon opportunities for growth [33] [88] . Every improvement in responsiveness or policy accuracy is a testament to the effectiveness of the safety foundation, not a reason to weaken it. This disciplined approach ensures that as the system grows more capable, it simultaneously grows more secure and more aligned with the user's fundamental principles of autonomy and self-determination.

Reference

1. Report of the International Bioethics Committee ... https://unesdoc.unesco.org/ark:/48223/pf0000378724
2. Towards a Governance Framework for Brain Data | Neuroethics https://link.springer.com/article/10.1007/s12152-022-09498-8
3. Neuroethics guidance documents: principles, analysis, and ... https://academic.oup.com/jlb/article/10/2/lsad025/7329870
4. Neurotechnologies under the Eye of Bioethics - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC9215686/
5. The neurotechnology patent landscape in a time of ... https://www.nature.com/articles/s41599-025-04757-4
6. Unveiling the neurotechnology landscape: scientific ... https://unesdoc.unesco.org/ark:/48223/pf0000386137
7. EEG-Based BCI Control Schemes for Lower-Limb Assistive ... https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2018.00312/full

8. Provable AI Ethics and Explainability in Medical ... https://www.mdpi.com/2079-9292/14/7/1294
9. Body of Knowledge - CS2023 https://csed.acm.org/wp-content/uploads/2024/01/Body-of-Knowledge-v1-bookmarksv2.pdf
10. Current Research Topics for the DEVCOM ARL BAA For ... https://defencescienceinstitute.com/wp-content/uploads/2024/06/AllTopicsDocumentforARLBAA7June2024.pdf
11. Computer Science https://arxiv.org/list/cs/new
12. 2022 roadmap on neuromorphic computing and engineering https://hal.science/hal-03872100v2/file/Christensen_2022.pdf
13. A Brain Foundation Model for Unified EEG and MEG Signals https://openreview.net/pdf/e1701399502f6010b41e08db60ccb6eda9d0746c.pdf
14. Impact of repeated exposure to CPM on CPM efficiency ... https://www.frontiersin.org/journals/pain-research/articles/10.3389/fpain.2025.1677563/full
15. Heterologous COVID-19 vaccine schedule with protein- ... https://www.sciencedirect.com/science/article/pii/S0163445325001707
16. Early-Stage Alzheimer's Disease Affects Fast But Not Slow ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11209650/
17. A Multimodal Framework for Prognostic Modelling of ... https://www.mdpi.com/2076-3417/16/2/763
18. A new guideline for predicting primary stability https://www.stanleyinstitute.com/blog/5-thread-guideline-a-new-guideline-for-predicting-primary-stability
19. Publication: Generative AI: Catalyst for Growth or Harbinger of ... https://openknowledge.worldbank.org/entities/publication/72ccdec4-3112-4b3f-9f07-25492da37c3f
20. energy, water, materials, and land limits of digital systems https://link.springer.com/article/10.1007/s44438-025-00018-8
21. Translating the One Security Framework for Global ... https://www.mdpi.com/2071-1050/18/2/1031
22. impact of generative artificial intelligence on socioeconomic ... https://academic.oup.com/pnasnexus/article/3/6/pgae191/7689236
23. MEDICAL PHYSICS in the BALTIC STATES 2023 https://www.researchgate.net/profile/Kirill-Skovorodko/publication/375838237_PATIENT_RADIATION_EXPOSURE_IN_PEDIATRIC_INTERVENTIONAL_RADIOLOGY_AT_VILNIUS_UNIVERSITY_HOSPITAL_SANTAROS_KLINIKOS/links/655f3651ce88b8703107c6b1/PATIENT-RADIATION-EXPOSURE-IN-PEDIATRIC-

INTERVENTIONAL-RADIOLOGY-AT-VILNIUS-UNIVERSITY-HOSPITAL-SANTAROS-KLINIKOS.pdf

24. Lakeshore owners attitudes toward their properties https://www.researchgate.net/publication/222519190_Sense_of_Place_as_an_attitude_Lakeshore_owners_attitudes_toward_their_properties
25. (PDF) Anthroposphere https://www.researchgate.net/publication/282268427_Anthroposphere
26. Application of GIS in Agriculture in Promoting Evidence- ... https://www.researchgate.net/publication/362656227_Application_of_GIS_in_Agriculture_in_Promoting_Evidence-Informed_Decision_Making_for_Improving_Agriculture_Sustainability_A_Systematic_Review
27. (PDF) Breastfeeding as a Fundamental Islamic Human Right https://www.researchgate.net/publication/333965705_Breastfeeding_as_a_Fundamental_Islamic_Human_Right
28. 2023 Index IEEE Transactions on Neural Networks and ... https://ieeexplore.ieee.org/iel7/5962385/10336252/10343530.pdf
29. Program in Chronological Order http://ieeexplore.ieee.org/iel5/4636107/4649055/04649058.pdf
30. JUNE 2023 https://www.ieee.org/ns/periodicals/NxtBooks/SP/PDF/SP_Jun2023.pdf
31. CRISPR and Artificial Intelligence in Neuroregeneration https://pmc.ncbi.nlm.nih.gov/articles/PMC12524701/
32. Speakers - AI for Good - ITU https://aiforgood.itu.int/summit25/speakers/
33. Gaussian Processes for Machine Learning (Adaptive ... https://dl.acm.org/doi/10.5555/1162254
34. Climate risk screening tools to enhance access to finance https://www.sciencedirect.com/science/article/pii/S2212096325000750
35. Pathways to a sustainable future: Exploring the synergy ... https://www.sciencedirect.com/science/article/pii/S2666188825007701
36. Reviewing studies of degrowth: Are claims matched by ... https://www.sciencedirect.com/science/article/pii/S0921800924002210
37. Ten new insights in climate science 2024 https://www.sciencedirect.com/science/article/pii/S2590332225001113
38. Conceptual Design Report for FAIR Computing https://arxiv.org/html/2511.01861v1
39. (PDF) The Future of Artificial Intelligence and ... https://www.arxiv.org/pdf/2509.02661v1

40. ARTIFICIAL INTELLIGENCE AND LIFE IN 2030 https://arxiv.org/pdf/2211.06318
41. LLMs4All: A Review of Large Language Models Across ... https://arxiv.org/html/2509.19580v5
42. AI Fairness in Practice https://arxiv.org/pdf/2403.14636?
43. Simulation-Based Inference: A Practical Guide https://arxiv.org/pdf/2508.12939
44. Conceptual Design Report for FAIR Computing https://arxiv.org/pdf/2511.01861
45. A theory of appropriateness with applications to generative ... https://arxiv.org/html/2412.19010v1
46. Contents https://arxiv.org/html/2509.02661v1
47. NeuroAI and Beyond https://arxiv.org/pdf/2601.19955
48. Ethics of neurotechnology: UNESCO adopts the first global ... https://www.unesco.org/en/articles/ethics-neurotechnology-unesco-adopts-first-global-standard-cutting-edge-technology
49. Ethics of neurotechnology https://www.unesco.org/en/ethics-neurotech
50. Draft text of the Recommendation on the Ethics ... https://unesdoc.unesco.org/ark:/48223/pf0000393395
51. Towards a draft text of a Recommendation on the Ethics ... https://unesdoc.unesco.org/ark:/48223/pf0000389438
52. First draft of the Recommendation on the Ethics ... https://unesdoc.unesco.org/ark:/48223/pf0000391444
53. Advancing neurotechnology while protecting the human ... https://www.unesco.org/en/articles/advancing-neurotechnology-while-protecting-human-brain-unesco-global-ethical-framework
54. First draft of a Recommendation on the Ethics ... https://unesdoc.unesco.org/ark:/48223/pf0000391074
55. Ethical issues of neurotechnology: report, adopted in ... https://unesdoc.unesco.org/ark:/48223/pf0000383559.locale=en
56. Liste des publications du Think Tank du PE https://www.europarl.europa.eu/thinktank/fr/research/advanced-search/pdf?keywords=64
57. List of publications from the EP Think Tank - European Union https://www.europarl.europa.eu/thinktank/en/research/advanced-search/pdf?keywords=64
58. Liste over publikationer fra EuropaParlamentets Think Tank https://www.europarl.europa.eu/thinktank/da/research/advanced-search/pdf?selectedDocumentType=&fakeSelectedKeywords=&sort=DATE_DOCU&fakeSelectedAuthor=&previousResultSetDateStart=01-05-2000&dateEnd=&dateStart=&previousSubmitter=WORD&selectedKeyword=&documentTypes=&page=2&selectedAuthor=

&selectedPolicyArea=&fakeSelectedPolicyArea=&previousResultSetDateEnd=11-06-2018&word=foreign%20fighters&facetAction=

59. Elenco delle pubblicazioni del Think Tank del PE https://www.europarl.europa.eu/thinktank/it/research/advanced-search/pdf?keywords=64
60. Neuro-Inspired Computing Using Resistive Synaptic ... https://www.scribd.com/document/848668056/Neuro-Inspired-Computing-Using-Resistive-Synaptic-Devices-Shimeng-Yu-Eds-Z-Library
61. Final report on the draft Recommendation on the Ethics of ... https://unesdoc.unesco.org/ark:/48223/pf0000393266
62. Draft Recommendation on the Ethics of Neurotechnology https://unesdoc.unesco.org/ark:/48223/pf0000394866
63. Loss of nuclear envelope integrity? No probLEM-BAF has it ... https://pubmed.ncbi.nlm.nih.gov/31201264/
64. The Future of Artificial Intelligence and the Mathematical ... https://www.arxiv.org/pdf/2509.02661
65. sheet1 https://clas.cas.cn/Y2021xwdt/Y2021tzgg/202310/P020231009340532191129.xlsx
66. Artificial Intelligence & Peace: A Living-Systems Approach ... https://www.academia.edu/122319955/Artificial_Intelligence_and_Peace_A_Living_Systems_Approach_to_AI_Ethics
67. Value Tensions in the Tokenization of Future Human-AI ... https://dl.acm.org/doi/fullHtml/10.1145/3663384.3663390
68. arXiv:2212.08967v1 [cs.AI] 17 Dec 2022 https://arxiv.org/pdf/2212.08967
69. A theory of appropriateness with applications to generative ... https://arxiv.org/pdf/2412.19010
70. Review on Sustainable Forestry with Artificial Intelligence https://dl.acm.org/doi/10.1145/3759259
71. FROM REPRESENTATION TO MEDIATION https://arxiv.org/pdf/2506.18743
72. The challenge of uncertainty quantification of large ... https://arxiv.org/html/2504.05278v1
73. Value Tensions in the Tokenization of Future Human-AI ... https://dl.acm.org/doi/pdf/10.1145/3663384.3663390
74. (PDF) Might artificial intelligence become part of the person ... https://www.researchgate.net/publication/365254909_Might_artificial_intelligence_become_part_of_the_person_and_what_are_the_key_ethical_and_legal_implications

75. Artificial Intelligence as a Disruptive Technology - Springer Link https://link.springer.com/content/pdf/10.1007/978-3-032-02827-3.pdf
76. List of publications from the EP Think Tank - European Union https://www.europarl.europa.eu/thinktank/en/research/advanced-search/pdf?keywords=2841
77. Linknovate | Profile for Columbia University https://www.linknovate.com/affiliation/columbia-university-210/all/?query=ultimately%20achieving%20unprecedented%20control
78. unesco https://research.bjmu.edu.cn/docs/2024-10/af391e2749cf4794b652945bebbb5b91.pdf
79. Final report on the draft text of the Recommendation ... https://unesdoc.unesco.org/ark:/48223/pf0000393400
80. Outcome document of the first meeting of the AHEG https://unesdoc.unesco.org/ark:/48223/pf0000389768
81. EP Ideju laboratorijas publikāciju saraksts - European Parliament https://www.europarl.europa.eu/thinktank/lv/research/advanced-search/pdf?policyAreas=SECDEF
82. Arxiv今日论文| 2025-11-27 http://lonepatient.top/2025/11/27/arxiv_papers_2025-11-27
83. Quantum Semantic Physics From QSIT Theory to Empirical ... https://www.researchgate.net/publication/397755773_Quantum_Semantic_Physics_From_QSIT_Theory_to_Empirical_Validation_of_the_Resonant_Mode_in_Multi-Agent_Systems
84. A Review of AI-Driven Quantitative Remote Sensing ... https://arxiv.org/html/2507.09081v1
85. Physics-Informed Spiking Neural Networks via ... https://arxiv.org/html/2511.21784v1
86. Mixture of Experts (MoE): A Big Data Perspective https://arxiv.org/html/2501.16352v1
87. Reusability of model components for environmental ... https://dl.acm.org/doi/10.5555/2781922.2782499
88. Hybrid EEG–fNIRS-Based Eight-Command Decoding for BCI https://pmc.ncbi.nlm.nih.gov/articles/PMC5314821/
89. fNIRS-EEG BCIs for Motor Rehabilitation: A Review - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC10740927/
90. EEG-fNIRS-based hybrid image construction and ... https://pmc.ncbi.nlm.nih.gov/articles/PMC9472125/

91. Wearable, Integrated EEG–fNIRS Technologies: A Review https://pmc.ncbi.nlm.nih.gov/articles/PMC8469799/
92. Explainable Artificial Intelligence Applications in Cyber ... https://arxiv.org/pdf/2208.14937
93. Review on Sustainable Forestry with Artificial Intelligence https://dl.acm.org/doi/full/10.1145/3759259
94. Computer Science https://www.arxiv.org/list/cs/new?skip=275&show=2000
95. Reinforcement Learning for Quantum Technology https://arxiv.org/html/2601.18953v1
96. Climate Data Practices: A Research Approach for HCI and ... https://dl.acm.org/doi/10.1145/3719346
97. Abstract https://arxiv.org/html/2508.07423v2
98. A review of empirical practices across implementation scales https://www.sciencedirect.com/science/article/pii/S0959652624032657
99. Operationalising AI governance through ethics-based auditing https://pmc.ncbi.nlm.nih.gov/articles/PMC9152664/
100. Empowering biomedical discovery with AI agents https://www.cell.com/cell/fulltext/S0092-8674(24)01070-5
101. Synthesis of Two Decades of US EPA's Ecosystem ... https://www.mdpi.com/2071-1050/13/15/8249
102. Ethical considerations for the use of brain–computer ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11542783/
103. Neural vs Neuromorphic Interfaces: Where Are We Standing? https://pubs.acs.org/doi/10.1021/acs.chemrev.4c00862
104. A Unified Frontier in Neuroscience, Artificial Intelligence ... https://arxiv.org/html/2507.10722v1
105. A blockchain security module for brain-computer interface ... https://www.sciencedirect.com/science/article/pii/S2772528621000303
106. AI Standards for Global Impact: From Governance to Action https://www.itu.int/dms_pub/itu-t/opb/ai4g/T-AI4G-AI4GOOD-2025-4-PDF-E.pdf
107. Neuralink's brain-computer interfaces: medical innovations ... https://www.frontiersin.org/journals/human-dynamics/articles/10.3389/fhumd.2025.1553905/full
108. Issue 12 - Volume 100 - Physica Scripta https://iopscience.iop.org/issue/1402-4896/100/12

109. Arxiv今日论文 | 2026-01-09 http://lonepatient.top/2026/01/09/arxiv_papers_2026-01-09.html